

The Reactive Literature

Measuring How Banking-Law Scholarship Reaches Public Policy, 2000–2025

Jacob Gerszten · Independent researcher *Working paper — draft of 2026-06-20*

Methods note: this study was conducted with AI assistance (Claude) as a research and drafting partner. The author set the questions and theses, selected and verified the sources and methods, and is responsible for the content. All data were assembled from open, terms-of-service-respecting sources; the pipeline is reproducible.

Abstract

Does legal scholarship on bank regulation influence policy, or merely comment on it? This paper assembles an original corpus of **1,915 banking-law review articles (2000–2025)** and links it to the federal policy record through free government data, in order to measure influence rather than assume it. The answer is stratified. In the aggregate, the literature is **reactive**: scholarly output tracks policy events at a two-to-three-year lag, and the topics that dominate the field — Dodd-Frank, crisis programs, Basel III — surge *after* the events that produced them, not before. But a small, identifiable **vanguard** sits upstream of policy, and its influence is visible across three independent channels: citation of specific articles in congressional and judicial documents, recurring expert testimony, and movement into government office. The paper documents both the reactive baseline and the upstream exceptions, develops a reproducible and license-respecting method for the linkage, and is deliberate about what "influence" can and cannot be shown. A companion interactive dashboard and a Bluebook-formatted citation dataset accompany the study.

1. Introduction

The proposition that scholarship shapes law is foundational to the legal academy and notoriously difficult to demonstrate. Citation studies measure attention within the academy; they rarely reach the question that matters for public law — whether ideas produced in law reviews actually move policy. Banking law is an unusually good place to ask the question. The field has a dense administrative apparatus (the Federal Reserve, OCC, FDIC, CFPB, Treasury, and FSOC), a well-documented revolving door between faculties and agencies, and a legislative cycle driven by crises that leave sharp temporal markers. If scholarly influence on policy is measurable anywhere, it is measurable here.

This paper distinguishes two questions that the literature often conflates. The first is a question of **supply**: what does the field produce — by whom, when, and on what topics? The second is a question of **influence**: does that production shape policy? Most quantitative work answers the first and infers the second. This study builds a policy-side evidence layer so the second can be tested directly.

The central finding is that the field, taken as a whole, **metabolizes** policy more than it **originates** it: banking-law scholarship is predominantly a reactive, interpretive literature that explains crises and

statutes after they land. But that aggregate masks a stratified structure. A small number of scholars — nameable, and identifiable across multiple independent channels — are genuinely upstream of policy. The contribution is therefore both substantive (a stratified account of influence that neither over- nor under-claims) and methodological (a reproducible, terms-of-service-respecting way to link a scholarly corpus to the federal policy record using only open data).

The paper proceeds as follows. Part 2 describes the corpus and how it was built, including a journal-of-record normalization that corrects a systematic mismeasurement of where scholarship is published. Part 3 maps the shape of the field. Part 4 tests whether scholarship leads or lags policy. Part 5 traces influence through three channels — citation, testimony, and personnel. Part 6 presents four documented case studies of the upstream vanguard. Part 7 synthesizes; Part 8 states the limitations plainly; Part 9 concludes.

2. The Corpus: Construction, Normalization, and Coverage

2.1 Construction

The corpus comprises **1,915 banking-relevant articles** published 2000–2025, filtered from roughly 8,400 deduplicated law-review records drawn from some 384 sources. The records were assembled exclusively from **open, license-respecting sources**: law-school institutional repositories (bepress/Digital Commons), open journal systems, publisher sitemaps, and the open bibliographic aggregators OpenAlex and Crossref. Sources behind restrictive terms — HeinOnline, Westlaw, Lexis, Google Scholar — were deliberately excluded, and machine-collection bans and crawl-delays (for example, those of the *Yale Law Journal* and *Stanford Law Review* websites) were honored rather than circumvented. This constraint is a feature, not merely a scruple: it makes the pipeline reproducible by anyone, and it forces the coverage analysis in Part 2.3 to be explicit about what open data can and cannot reach.

Each article was classified into a 30-category doctrinal taxonomy (capital and Basel, resolution and "too big to fail," consumer financial protection, fintech and digital assets, Federal Reserve governance, and so on) by a rule-based classifier with hand-audited edge cases. The taxonomy deliberately excludes adjacent fields — corporate, securities, and bankruptcy law — except where a piece bears directly on bank regulation. An author layer links 1,477 distinct authors to affiliations, yielding a ranked list of the field's most prolific scholars and a co-authorship network.

2.2 Journal-of-Record Normalization and Bluebook Citations

A subtle measurement problem emerged and was corrected. Many articles enter the corpus through a law school's faculty repository or as a self-archived working paper, so their recorded venue is the *repository* ("Michigan Faculty Scholarship") or "SSRN Electronic Journal" rather than the journal in which the piece actually appeared. Counting publication venues from those labels systematically **undercounts** elite reviews. The clearest illustration: Rauterberg and Zhang's *Shadow Banking and Securities Law* is filed in the corpus under a Michigan repository but was published in the **Stanford Law Review** — 77 Stan. L. Rev. 563 (2025). Neither OpenAlex nor Crossref had yet linked that 2025 article to Stanford's ISSN; the corpus captured it only because the author self-archived it.

To fix this, a journal-of-record pass parses each repository landing page's citation metadata and re-attributes the piece to its journal of publication. The result feeds a **Bluebook citation dataset**: every record carries a citation built from its journal of record — never SSRN or a repository. Of 1,915 records, **1,326 (69%) have a complete citation** (author, title, volume, journal, first page, year); 341 have a known journal but pending volume/page; and 248 (13%) are flagged for manual completion because their published venue cannot be recovered from open sources alone. Flagging the residual, rather than mislabeling it, keeps the dataset honest.

2.3 Coverage, Stated Honestly

The honest coverage picture is better than venue labels suggest, with one genuine and instructive exception. After journal-of-record normalization, banking-relevant coverage is solid across every flagship review — California (18), Chicago (17), Vanderbilt (12), Fordham (11), Harvard (10), Cornell (10), Columbia (9), Penn (9), Yale (6) — and **twelve flagship pieces were recovered** from repository attribution. The *Yale Law Journal*, long assumed a gap because its website bars automated collection, turns out to be well covered through faculty self- archiving and open aggregators. The *Stanford Law Review*, by contrast, contributes essentially one banking-regulation article in the entire window. That is **not** a harvesting failure: a direct check of all 325 Stanford pieces in OpenAlex confirms the journal simply publishes little bank-*regulation* scholarship, its financial-law output running to corporate and bankruptcy topics that fall outside the taxonomy. The genuine, irreducible undercounts are the paywalled-only venues — most importantly the *Banking Law Journal* (LexisNexis) — which open data cannot reach.

3. The Shape of the Field

Annual output roughly tripled across the period, from the high-20s/low-40s in the early 2000s to a post-crisis plateau above 100, peaking at 121 articles in 2012. The composition tells the field's story. **Dodd-Frank-specific scholarship** is the signature: essentially nonexistent before 2010, it explodes to 139 articles in the 2010–2014 window before receding. **Fintech, crypto, and digital-asset** scholarship is the secular growth story, rising from a trickle to the single largest topic by the 2020s. **Consumer financial protection** is a steady heavyweight throughout. Crisis programs, "too big to fail," and Basel-style capital all spike around 2008–2014 and then decline. The shape of production already hints at the paper's thesis: the field's largest movements are organized around external shocks and the statutes that answer them.

4. Does Scholarship Lead or Lag Policy?

To test the lead/lag question directly, the topic-by-year series were aligned against a hand-built timeline of 34 major banking-policy events (statutes, marquee rules, crises, and executive actions), each mapped to the corpus topic it most implicates. Two analyses follow.

First, an event study. For each event mapped to a topic with sufficient signal, the topic's output in the three years before the event is compared with the event year plus the three years after. **Twelve of seventeen** reliably-windowed events show the topic surging *after* the event, not before. The cleanest cases are topics effectively *created* by an event — Dodd-Frank, the crisis-era emergency programs, the Basel III transposition — which barely exist beforehand and explode afterward. The apparent "anticipatory" cases are not counterexamples; they are continuously-studied plateau topics (consumer protection, securitization) where a single event produces no visible surge.

Second, a cross-correlation between total banking output and a policy-event impulse series peaks at a **lag of +2 to +3 years ($r \approx 0.43$)**: scholarly output correlates most strongly with the policy events of two to three years earlier. Both analyses point the same way. In the aggregate, the literature follows policy.

This is the paper's baseline, and it should be read with discipline. Co-movement is not causation, and a positive lag is consistent with the *reverse* causal story — that policy debates generate scholarship. That is precisely the point: the aggregate evidence supports a reactive, interpretive account of the field, against which the upstream exceptions in Parts 5 and 6 stand out.

5. Tracing Influence: Three Channels

If most scholarship is reactive, where is the influence? The study traces it through three independent channels linked to the federal record via free government APIs (GovInfo, which spans congressional hearings and reports, the Federal Register and CFR, and federal court opinions; and Congress.gov). Each channel has a distinct mechanism, and the paper is careful to distinguish them.

5.1 Citation in Federal Documents

Articles were matched to federal documents by an **exact-title phrase search confirmed by author-surname co-occurrence** — a high-precision rule that eliminates generic-phrase false positives (the title "The Too Big to Fail Problem" alone matches scores of hearings; requiring "Omarova" in the same document collapses that to the real citations). The method is precise but recall-limited: it captures verbatim-title references, not citations that use only a Bluebook abbreviation, so the counts are **lower bounds**. (Two recall corrections — handling parenthetical subtitles and short distinctive titles — raised the confirmed count from 105 to 152.)

Of 1,764 distinctively-titled articles, **152 (≈9%) are cited by exact title in at least one federal document**, across confidence tiers: roughly 74 with a congressional citation (the high-confidence core — a title plus author in a hearing or report is almost always a real reference), 62 with a federal-court citation (noisier, given common surnames), and 16 in the Federal Register or CFR. Confirmed examples include Frank Partnoy's credit-rating-agency scholarship appearing in the very hearing convened on rating-agency regulation and in the Senate's *Wall Street and the Financial Crisis* report; Mehrsa Baradaran's *How the Poor Got Cut Out of Banking* in Community Reinvestment Act and central-bank-digital-currency hearings; and the Ricks–Menand "FedAccounts"/public-option work in the Senate Banking CBDC hearing. The honest caveat: a citation evidences awareness and legitimation, not that an outcome changed — and a hearing citation is sometimes the scholar entering their own work into the record, a genuine but weaker form of uptake than third-party citation.

5.2 Congressional Presence

A complementary and stronger signal searches the congressional record for each scholar's full name (date-filtered to 2000–2026 to exclude historical namesakes). **115 of 223 corpus scholars appear in at least one congressional hearing or committee report**. The leaders, after flagging common-name collisions, are Michael Barr (67 documents), Daniel Tarullo (53), Adam Levitin (48), John Coffee (38), Margot Saunders (32), Susan Wachter (25), David Skeel (25), Raphael Bostic (25), Frank Partnoy (18), and Chris Brummer (17). This is the strongest quantitative evidence of the literature entering the policy conversation — though, again, "presence" includes testifying, being cited by another witness, and being named by a member, which are different mechanisms.

5.3 The Revolving Door

The cleanest channel is personnel: ideas travel with people. Cross-referencing the scholar list against government service yields a clear hierarchy. **Senate-confirmed principals**: Barr and Tarullo (Federal Reserve Governors). **Failed confirmations** — themselves evidence of the pipeline and its political friction: Saule Omarova (OCC) and Chris Brummer (CFTC). **Non-confirmed staff and appointees**, the dominant mode: Warren and Levitin (TARP Congressional Oversight Panel), Menand, Ricks, McCoy, and Gelpern (Treasury and agencies), Gerding (SEC division director), and Skeel (PROMESA board). **Advisors**: Judge (Treasury's OFR committee), Wilmarth and Allen (Financial Crisis Inquiry Commission), Hockett (FRBNY), and Baradaran (transition team).

6. The Upstream Vanguard: Case Studies

Four documented chains show influence running from scholarship to policy, each labeled by its operative mechanism.

(a) The "money view" and central-bank digital currency. A scholarly strand — Morgan Ricks's *Money Problem*, the Hockett–Omarova "finance franchise," and the Ricks–Menand work on FedAccounts and a public option for bank accounts — is visible in the digital-dollar legislative proposals, the Federal Reserve's CBDC research, and the contested Omarova nomination to the OCC, which made the academic argument politically salient. Mechanism: principally testimony and discourse rather than third-party citation.

(b) "Unsafe at Any Rate" and the CFPB. The field's canonical scholarship-to-institution chain: Elizabeth Warren's 2007 essay proposing a financial-product safety commission became, through Dodd-Frank's Title X, the Consumer Financial Protection Bureau — which Warren then stood up. Adam Levitin's recurring testimony and Patricia McCoy's move into the Bureau (where her subprime-lending scholarship informed mortgage supervision) extend the chain; McCoy is the cleanest scholarship-to-rule link.

(c) Financial inclusion and postal banking. Mehrsa Baradaran's work on financial exclusion and the racial wealth gap traces to the Postal Banking Act, the 2021 USPS banking pilot, and racial-wealth-gap hearings — an unusually upstream scholar, though several of the apparent citations are her own testimony.

(d) The scholar-regulators. Supervision and resolution policy carried by scholars-in-office: Barr (Treasury architect of Dodd-Frank, then the Fed's Vice Chair for Supervision), Tarullo (architect of the post-crisis stress-testing regime), Kress (whose bank-merger scholarship fed the 2024 revision of merger review), and Skeel and Wilmarth on resolution and "too big to fail." This is the cleanest personnel channel — with the honest qualification that for Tarullo the causation often runs policy → literature, as he wrote influential critiques *after* shaping the rules.

7. Synthesis: A Stratified Account of Influence

Reading the channels together, the scholars who register on *several* at once are the genuine upstream nodes — Barr (67 hearings, cited, Senate-confirmed), Levitin (48 hearings, served), Wachter (25 hearings, six citations), Partnoy (18 hearings, seven citations), Omarova (citations, testimony, a failed nomination), Skeel (testimony, PROMESA). The picture that emerges is not "scholarship drives policy" but a **stratified influence structure**: a reactive literature in the aggregate, atop which a small, nameable vanguard testifies, is cited, and serves. The field-mapping and the influence analysis reinforce each other: the topics that surge after shocks are the field's bulk, while the upstream exceptions are concentrated in a handful of scholars whom the network and personnel data independently identify.

8. Limitations

The study is built to be honest about its limits, and they are real.

- **Timing is not causation.** The aggregate lag is equally consistent with policy debates generating scholarship; the paper treats the reactive finding as the null, not as proof of one-directional causation.
- **Citation and presence measure awareness and legitimization**, not that an outcome changed.
- **Self-testimony versus third-party citation.** Many congressional appearances are the scholar testifying — genuine entry into the debate, but a weaker claim than independent uptake.
- **Recall is bounded.** Exact-title matching misses abbreviation-only citations, so counts are lower bounds; name-matching is noisy and collisions are flagged, not hidden.
- **Informal channels are invisible.** Advising, amicus briefs, op-eds, and practitioner blogs — plausibly large conduits — leave no trace in this data.
- **Recent years undercount** because of publication lag; the partial final year is dropped.
- **Provenance gaps.** About 13% of records lack a journal of record recoverable from open sources and are flagged rather than guessed.

9. Conclusion

Banking-law scholarship is, in the main, a reactive and interpretive literature: it shapes how policy is understood and implemented far more than it sets the agenda. But a stratified reading rescues the influence thesis from both overstatement and dismissal. A small vanguard — identifiable, convergently, across citation, testimony, and personnel — is genuinely upstream of policy, and its members can be named. The methodological payoff is equally concrete: a reproducible, license-respecting way to study scholarly influence on public law using only open government data, together with an honest ledger of where the evidence is strong (congressional citation and presence; the revolving door) and where it is not (verbatim-title recall; informal channels). The same infrastructure that maps the field can, with care, measure its reach.

Appendix: Data and Reproducibility

- **Corpus & classification:** ~1,915 banking-relevant articles (2000–2025), 30-topic taxonomy, 1,477 authors; open, ToS-respecting sources only.
- **Citation dataset (Bluebook):** `banking_citations.csv` — 1,326 full / 341 journal-only / 248 flagged for manual completion; journal of record never SSRN or a repository.
- **Policy linkage:** GovInfo (congressional, judicial via USCOURTS, FR/CFR), Congress.gov.
- **Key outputs:** `leadlag_summary.md`, `citation_ledger.{md,csv}`, `congressional_presence.{md,jsonl}`, `revolving_door.md`, `influence_scorecard.{md,csv}`, `flagship_coverage.md`, `case_studies.md`, `policy_event_timeline.csv`.

- **Interactive companion:** the public dashboard at <https://gersztenj.github.io/banking-law-scholarship-map/>.
- **Honesty conventions:** no paywalled/ToS-restricted scraping; crawl-delays honored; name-match collisions and recall limits disclosed in situ.